

How to Root 10^7 Phones with One Exploit

James Fang **KEEN TEAM**

About Me

- Root guy of Keen Team
 - With support from the whole team
- Responsible for many PHA (root) application
 - Pingpong Root
 - Mate 7 Root
 - KingRoot and SDK (with KR team)

Why Root?

- Fun
 - Xposed, Greenify, etc
 - New ideas from community (XDA, Gfan)
 - Full control over your device
- Secured
 - Root manger
 - Additional exploit mitigation

Android Root (History)


- 2008: T-Mobile G1 hidden root console
 - Run terminal emulator app
 - Launches telnetd as root

```
Trying 192.168.0.88...
Connected to 192.168.0.88.
Escape character is '^]'.
# id
uid=0(root) gid=0(root)
# cd /proc
# cat version
Linux version 2.6.25-01843-gfea26b0 (android-build@apa27.mtv.corp.google.com)
(gcc version 4.2.1) #6 PREEMPT Mon Oct 6 14:13:36 PDT 2008
```

- Feature vs. bug



Android Root (History)



Android Root (History)

- OTA
 - Volez by Zinx
 - Improper parsing of signed update zip file
 - Master Key vulnerability “family”
 - Cydia impactor
 - system->root



Android Root (History)

- Daemon
 - setuid failure
 - Shared memroy
 - Memory corruption
 - GingerBreak
 - zergRush
 - File perms & symlinks ☺
 - Many were there in init.rc



Android Root (History)

- Kernel
 - Wunderbar/asroot
 - FramaRoot by alphazain
 - /dev/exynosmem
 - /dev/exynosmem (again :D)
 - CVE-2013-6282
 - towelroot
 - ...




How to ~~Cook~~ Root



Compounds

- Vulnerability & Exploit
 - Attack surface
 - In kernel/privileged process (to reach kernel)
- Mitigation Bypass
 - SELinux
 - Root mitigation (Samsung)
 - (NAND)? write protection

Challenges



Challenges

三星Galaxy Note II				三星Galaxy S IV				三星Galaxy S III																
三星Galaxy Note III				三星Galaxy Grand 2		三星Galaxy S V		三星GT-S7562		三星GT-S7568		三星Galaxy Ace		三星GT-S7568										
				三星Galaxy Win	三星Galaxy Tr		三星Galaxy		三星SM-G5		三星GT-S		三星Galax		三星GT-i9		三星GT-S		三星R830					
					三星Galaxy G	三星A5000		三星Gala		三星Gala		三星n910		三星SCH		三星i879		三星SM		三星Gal				
						三星Galaxy		三星 GAL		三星SM		三星GT		三星SC		三星Ga		三星GT		三星Ga		三星G		
				三星Galaxy S	三星GT-i8262		三星Galaxy		三星 GAL		三星SM		三星GT		三星SC		三星Ga		三星GT		三星Ga		三星G	
					三星i8268	三星Galaxy		三星Gala		三星G5		三星I91		三星GT		三星Ga		三星N9		三星n7		三星Ga		
						三星SM-3		三星GT		三星G		三星G		三星N		三星C		三星C		三星C		三星C		
				三星Galaxy Note	三星i739		三星Galaxy		三星GT-S7		三星 Ga		三星I6		三星S		三星G		三星		三星		三星	
					三星SM-G381		三星SCH-I8		三星GT-i		三星G3		三星W		三星G		三星G		三星I		三星		三星	
					三星SM-G381		三星Galaxy		三星SM-G		三星Gal		三星N9		三星G		三星G		三星N		三星C		三星	
三星Galaxy S II				三星Galaxy Mega 5.8		三星Galaxy		三星SM-G		三星Gal		三星N9		三星G		三星G		三星N		三星C				
				三星Galaxy		三星SM-G		三星Gal		三星N9		三星G		三星G		三星N		三星C		三星				
				三星Galaxy		三星SM-G		三星Gal		三星N9		三星G		三星G		三星N		三星C		三星				

Challenges

- Manufacturers and models
 - Kernel behaviors (syscall?)
 - Root mitigation
- Same model, different ROM
 - Kernel symbol locations
 - Structures, offsets
- Device database?
 - Not the best solution



A Case Study (Early 2014)

- CVE-2013-6282
 - put_user/get_user
 - Lack of boundary check
- Read & write
 - ptrace->put_user->write anywhere
 - setsockopt->get_user->read anywhere
- Targeting all ARMv6+ (by getsockopt)

A Case Study (Early 2014)

- “Perfect vulnerability”
- Thinking of these problems:
 - Achieve root with read & write
 - Achieve root with write only
 - Without device database
 - SELinux (Nexus and Samsung)
- “Write-only” bugs are quite common


DKOM

- Direct Kernel Object Manipulation
- Straight forward when we have R/W
- Our goal
 - Get root uid/gid
 - Get full capabilities
 - Get u:r:init:s0
 - Patch SELinux policy
 - Patch/bypass root mitigations
- Where are the “objects”?

DKOM

- Task information
 - struct thread_info => on kernel stack
 - struct task_struct => in heap
 - thread_info.task -> task_struct
- Leak SP (towelroot)
- Search and manipulate
 - Useful when code execution is inconvenient

The diagram illustrates the memory layout of the kernel stack and the `task_struct` structure. On the left, the `kernel stack` is shown as a vertical bar with a green top section, a grey middle section, and an orange bottom section labeled `struct task_struct *task;`. The stack grows from high memory addresses at the top to low memory addresses at the bottom, with a `+0xc` offset indicated. To the right, the `task_struct` structure is shown as a vertical bar with a light blue top section and a light blue bottom section. The bottom section contains `struct mm_struct *mm;`, `struct plist_node pushable_tasks;`, `struct list_head tasks;`, and `void *stack;`. The `tasks` field is marked with `+???` and the `stack` field with `+4`. A red line with arrows shows the flow of pointers: from the `tasks` field in `task_struct` to the `init_task` structure (a red bar with a red pushpin), and from the `stack` field in `task_struct` to the `kernel stack`. Other blue bars represent other `task_struct` instances in the system, connected by red arrows.

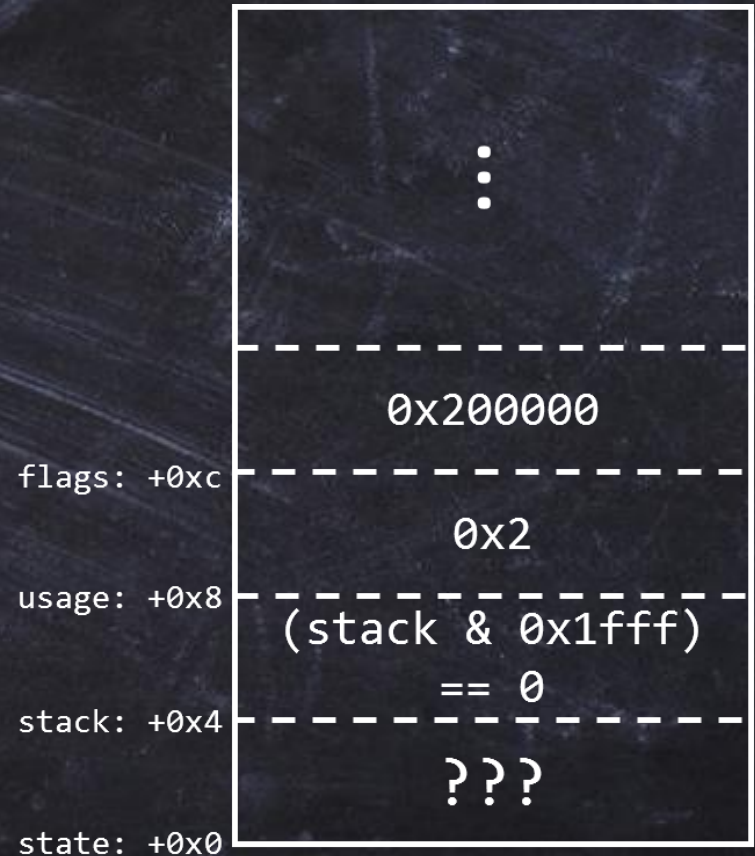


DKOM

- `init_task` leads to the doubly-linked list
- Usually a symbol in `kallsyms`
 - We don't have root yet
- Any pattern?
- In `/arch/arm/kernel/init_task.c`
 - `struct task_struct init_task = INIT_TASK(init_task);`
(take ARM32 as an example)


DKOM

- Always starts with this pattern:
 - .state => unknown
 - .stack => 2-page (8K) aligned address
 - .usage => 0x2
 - .flags => 0x2000000
- Search range?
 - /proc/iomem



DKOM

- Next: offset of tasks
- Structure definition:



```
struct task_struct {  
    volatile long state;  
    void *stack;  
    ...  
    struct list_head tasks;  
#ifdef CONFIG_SMP  
    struct plist_node pushable_tasks;  
#endif
```

- For all modern multi-core phones, CONFIG_SMP is always 'y'

DKOM

- Check `init_task.h`, `list.h` and `plist.h` in `/include/linux`
- `tasks + pushable_tasks` looks like:
 - Two kernel pointers
 - `0x00000008c`
 - Two kernel pointers
- Good anchor once we located `init_task`

DKOM

- How to identify my process?
 - Random string =>
char comm[TASK_COMM_LEN]
 - Anchor to identify task
 - It will also find task->cred for you 😊

```
const struct cred __rcu *cred;  
struct cred *replacement_session_keyring;  
char comm[TASK_COMM_LEN];
```

DKOM

- So now we have:
 - `init_task`
 - Traverse task list for any task
 - Identify task by comm
 - Offset of `task->cred`
- `struct cred` is simple and easy?
 - For uid/gid
 - How about `u:r:init:so?`

DKOM

- `sid <=> SELinux context`
- `1 <=> u:r:kernel:so`
- `?? <=> u:r:init:so`
- Two approaches
 - `untrusted/shell => kernel => init`
 - Policy may block it
 - Set `sid` directly
- Look for “init” task 😊

-*> MEMORY DUMP <*-

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
	01	00	00	00	<u>30</u>	<u>00</u>	<u>00</u>	<u>00</u>	00	00	00	00	00	00	00	00
	00	00	00	00	00	00	00	00								

-WX-WX-WX

- But what if only write anywhere is possible?
- Real example:
 - Broadcom /dev/uio1
 - ioctl: Write a string to given address
 - How to exploit?
- Overwrite syscall table and run shell code?

-WX-WX-WX

- Interesting entry in process map:

```
shell@test:/ $ cat /proc/self/maps
40000000-40020000 r--s 00000000 00:0b 4103      /dev/__properties__
40031000-40033000 rw-p 00000000 00:00 0
4006d000-40092000 r-xp 00000000 b3:1c 353      /system/bin/mksh
40092000-40093000 r--p 00000000 00:00 0
.
.
.
401c1000-401c4000 rw-p 0004a000 b3:1c 1704      /system/lib/libc.so
401c4000-401d2000 rw-p 00000000 00:00 0
4030e000-40319000 rw-p 00000000 00:00 0      [heap]
bef49000-bef6a000 rw-p 00000000 00:00 0      [stack]
ffff0000-ffff1000 r-xp 00000000 00:00 0      [vectors]
```

- What's inside?

-WX-WX-WX

- In /arch/arm/kernel/entry-armv.S

```
.LCvswi:
    .word    vector_swi
...
__vectors_start:
ARM(      swi      SYS_ERROR0  )
THUMB(     svc      #0          )
THUMB(     nop          )
    W(b)      vector_und + stubs_offset
    W(ldr)     pc, .LCvswi + stubs_offset
    W(b)      vector_pabt + stubs_offset
...
    W(b)      vector_fiq + stubs_offset

    .globl    __vectors_end
__vectors_end:
```

- At 0xffff0008, location of vector_swi is leaked by an LDR instruction.
- loc_vector_swi =
*(0xffff0008 + 8 + ((*0xffff0008) & 0xfff))

-WX-WX-WX

- Symbol `sys_call_table` is always 3 symbols away from `vector_swi`

`vector_swi`

`__sys_trace`

`__sys_trace_return`

`__cr_alignment`

`=>sys_call_table`

- All defined in `/arch/arm/kernel/entry-common.S`
- `loc_sys_call_table = loc_vector_swi + len(4 following stubs)`

-WX-WX-WX


- Length of vector_swi depends on:
 - CONFIG_OABI_COMPAT - Allow old ABI binaries to run with this kernel (EXPERIMENTAL)
 - CONFIG_SECCOMP - Enable seccomp to safely compute untrusted bytecode
- In most cases both are set as N.
- If set:
 - CONFIG_SECCOMP => 6 instructions, 24 bytes
 - CONFIG_OABI_COMPAT => More complex, 28 bytes

-WX-WX-WX

- CONFIG_SECCOMP: PRCTL(PR_GET_SECCOMP, ...)
- CONFIG_OABI_COMPAT: sys_oabi_semop
- Algorithm is simple
 - Set BASE_OFFSET
 - If (CONFIG_SECCOMP) BASE_OFFSET += 24
 - If (CONFIG_OABI_COMPAT) BASE_OFFSET += 28
 - Align
 - loc_sys_call_table = loc_vector_swi + BASE_OFFSET
- Patch unused syscall entry and run shell code
- Assuming no PXN or RO kernel text
(True for most devices before 2014)

Conclusion

- A good vulnerability
 - Hardcode-free exploitation
 - Fine tune
- => Decent coverage



What's next?

- Root mitigation bypass
 - Samsung (STRICT_SEC)
 - Huawei fblock
 - Sony RIC
- CVE-2015-3636
 - How to root 10^7 phones again?
 - BH 2015

Ah! Universal Android Rooting is Back
Present by Wen Xu



Thank you